

Problem 1

If G is not connected, we can analyze each connected component independently, subtracting the combined sizes of all other components from k . Thus, for the remainder of the solution, we assume that G is connected.

If there exists a subset of vertices X such that $|X| \leq k$ and $G \setminus X$ is a tree, then $\text{tw}(G) \leq k + 1$, as $\text{tw}(G \setminus X) = 1$, where tw denotes treewidth. We will use the algorithm presented in the lecture, which runs in time $27^l \cdot l^{\mathcal{O}(1)} \cdot n^2$, where l is the target treewidth and n is the number of vertices. We apply it to G with $l = k + 1$. The algorithm will yield one of two possible outcomes:

1. Confirmation that $\text{tw}(G) > k + 1$.

In this case, we conclude that no such subset X exists.

2. A tree decomposition of width at most $4k + 8$.

The remainder of the solution focuses on this scenario.

We now proceed with dynamic programming, assuming the decomposition is nice. For any subtree H of the decomposition and its boundary ∂H , let $f : V(\partial H) \rightarrow \{0, \dots, 4k + 9\}$. Define $\text{dp}_H(f)$ as the size of the minimum set $Y \subseteq V(H \setminus \partial H)$ satisfying the following conditions:

1. $H \setminus (Y \cup f^{-1}(0))$ is a forest,
2. for all $v, u \in V(\partial H) \setminus f^{-1}(0)$, v and u are in the same connected component of $H \setminus (Y \cup f^{-1}(0))$ if and only if $f(v) = f(u)$.

If no such set Y exists for a given f , we define $\text{dp}_H(f) = \infty$.

To compute dp_H , we consider the following cases:

1. $H = \text{introduceVertex}(H', v)$

Here, $\text{dp}_H(f) = \min(\{\text{dp}_{H'}(f') : f = f'[v \mapsto f(v)] \wedge f'^{-1}(f(v)) = \emptyset\})$, where $\min(\emptyset) = \infty$. We use the notation $f[a \mapsto b]$ to denote the function g defined by:

$$g(x) = \begin{cases} b, & \text{if } x = a, \\ f(x), & \text{otherwise.} \end{cases}$$

2. $H = \text{introduceEdge}(H', v, u)$

In this case, $\text{dp}_H(f) = \min(\{\text{dp}_{H'}(f') : f'(v) \neq f'(u) \wedge (f'(w) = f'(v) \vee f'(w) = f'(u)) \Rightarrow f(w) = f(v) = f(u)\})$.

3. $H = \text{forgetVertex}(H', v)$

Here, $\text{dp}_H(f) = \min(\{\text{dp}_{H'}(f') + [k = 0] : k \in \{0, \dots, 4k + 9\} \wedge f' = f[v \mapsto k]\})$, where $[P]$ denotes the Iverson bracket, i.e., $[P] = 1$ if P is true, and $[P] = 0$ otherwise.

4. $H = \text{merge}(H', H'')$

In this case $\partial H' = \partial H''$, and we calculate dp_H as $\text{dp}_H(f) = \text{dp}_{H'}(f) + \text{dp}_{H''}(f)$.

The answer is $\text{dp}_T(\text{empty function}) \leq k$, where T is the entire decomposition.

The total complexity is bounded by

$$27^{k+1} \cdot (k+1)^{\mathcal{O}(1)} \cdot n^2 + n^{\mathcal{O}(1)} \cdot ((4k+10)^{4k+10})^2,$$

thus this algorithm is FPT when parameterized by k .

Problem 2

We apply the color-coding technique. Let $c = k(l-1)$ denote the maximum total number of vertices along the paths, excluding the start and finish vertices. We color each vertex (excluding the start and finish vertices) with one of c colors. For each $i \in \{1, \dots, k\}$, let $\text{dp}_i[v][C]$ be true if and only if there exists a path from s_i to v such that the set of vertex colors on this path equals exactly C , with no two vertices sharing the same color. Here, we only consider subsets C of size at most $l-1$.

The values of dp_i are initially set to false and are then computed using the following rules:

1. $\text{dp}_i[s_i][\emptyset] = \text{true}$.
2. $\text{dp}_i[v][C] = \bigvee_{(u,v) \in E(G)} (\text{color}[v] \in C \wedge \text{dp}_i[u][C \setminus \{\text{color}[v]\}])$, for all $v \in V(G) \setminus \bigcup_{j=1}^k \{s_j, t_j\}$.
3. $\text{dp}_i[t_i][C] = \bigvee_{(u,t_i) \in E(G)} \text{dp}_i[u][C]$.

To compute these values correctly, ensuring that no dp entry is referenced before it has been calculated, subsets C are considered in increasing order of size. Subsequently, $\text{dp}_i[v][C]$ is filled for all $v \in V(G)$.

The number of such subsets is bounded by:

$$\sum_{j=0}^{l-1} \binom{c}{j} \leq \sum_{j=0}^{l-1} c^j = \frac{c^l - 1}{c - 1} \leq c^l = (k(l-1))^l \leq (kl)^l.$$

Thus, this computation requires $\mathcal{O}((kl)^l k(n+m))$ time, where $n = |V(G)|$, and $m = |E(G)|$.

The number of k -tuples of such subsets is bounded by $((kl)^l)^k = (kl)^{kl}$. We iterate over these tuples, performing the following for each (C_1, \dots, C_k) :

1. Verify that the subsets C_1, \dots, C_k are pairwise disjoint.

This step can be done in $\mathcal{O}(k^3 l)$ time.

2. Check whether $\text{dp}_1[t_1][C_1] = \dots = \text{dp}_k[t_k][C_k] = \text{true}$.

In this case, the answer to the problem is „Yes.”.

The time complexity of one iteration is $\mathcal{O}((kl)^l k(n+m) + (kl)^{kl} k^3 l)$.

The probability of a single iteration failing to find a solution, assuming one exists, is $1 - \frac{c!}{c^c}$, as there are c^c ways to assign c colors to c vertices, and $c!$ of these avoid duplicate colors. Repeating the iteration e^c times reduces the failure probability to:

$$\left(1 - \frac{c!}{c^c}\right)^{e^c} < \left(1 - \frac{1}{e^c}\right)^{e^c} < 1 - \frac{1}{e},$$

a constant.

The total time complexity is:

$$\mathcal{O}(e^{kl}((kl)^l k(n+m) + (kl)^{kl} k^3 l)),$$

which is FPT when parameterized by $k+l$, as $kl \leq (k+l)^2$.

This algorithm can be determinized using an (n, c, c) -splitter, as discussed in the lecture. The splitter, with size $e^c c^{\mathcal{O}(\log c)}$, can be found in time $e^c c^{\mathcal{O}(\log c)} n \log n$. The resulting complexity is:

$$e^{kl} (kl)^{\mathcal{O}(\log kl)} (n \log n + \mathcal{O}((kl)^l k(n+m) + (kl)^{kl} k^3 l)),$$

which is also FPT when parameterized by $k+l$, completing the proof.